




An Introduction to Deep Learning

Drew Minkin
Lead Data Science Instructor
Jan 5 2023



Topics Overview

1. The Basics of Machine Learning
2. Machine Learning , Deep Learning & Artificial Intelligence
3. Starting with Deep Learning: Keras
4. TensorFlow vs PyTorch and Model Production
5. ONNX and Model Consumption





Who's Drew?

Previous Incarnations

SQL Support/Consulting 2000-06



Data Scientist 2010 – Now



5 Startups 2007-19



Analytics Architect/Dev



Entrepreneur



Current Incarnation

MCT, Data Science Instructor





The Basics of Machine Learning






Machine Learning , Deep Learning & Artificial Intelligence



Machine Learning Process

Model Production



Machine Learning Modeling Families



Estimation



Association



Clustering




Classification



Forecasting

Machine Learning Modeling Families




Machine Learning Modeling Families



- Hard and Soft Groupings
- Profiles of Subgroups
- Likenesses and Differences



Estimation



Association



Clustering



Classification



Forecasting

Machine Learning Modeling Families



Estimation



Association



Clustering

Predicting a Continuous Distribution
▪ Many Different Measures of Accuracy



Classification



Forecasting

Machine Learning Modeling Families



Estimation



Association



Clustering



Classification



Forecasting

- Input of measure over time and related series
- Predictions generated for short term trends
- Based on cycles and events

Machine Learning Modeling Families



Estimation



Association

- Collaborative Filtering
- Identify cross-sell
- Identify sequential, next-sale
- Make purchase recommendations
- Complex event associations



Clustering



Classification



Forecasting

What is Artificial Intelligence?


- Software that exhibits human-like capabilities, such as:



Visual Perception



Text Analysis



Speech




Decision Making

Machine Learning vs Deep Learning




Machine Learning



- A more complex **representation** of data
- Processing of **complex feature spaces**


Deep Learning



- **Feature Engineering baked into Models**
- **Deep** layering of algorithm, not wide.

<https://www.xenonstack.com/blog/static/public/uploads/media/machine-learning-vs-deep-learning.png>

Deep Learning in Machine Learning Algorithm Context



- LR, NN and other algorithms part of **hierarchical layers** of DL model
- **Transfer Learning**
Learning models on same type of data

Data Science, Machine Learning, and AI



Artificial Intelligence
Intelligent software apps and agents

Machine Learning
Use of data and algorithms to train predictive models

Data Science
Application of mathematical and statistical techniques to analyze data

Deep Learning, Machine Learning, and AI



Artificial Intelligence

Algorithms that mimic the intelligence of humans, able to resolve problems in ways we consider “smart”. From the simplest to most complex of the algorithms.

Machine Learning

Algorithms that parse data, learn from it, and then apply what they've learned to make informed decisions. They use human extracted features from data and improve with experience.

Deep Learning

Neural Network algorithms that learn the important features in data by themselves. Able to adapt themselves through repetitive training to uncover hidden patterns and insights.


Azure Deep Learning Platforms – Azure Only



AZURE
MACHINE LEARNING STUDIO



AZURE
SYNAPSE ANALYTICS



AZURE
COGNITIVE SERVICES

Azure Deep Learning Platforms - Databricks



AZURE
DATABRICKS

- Git Support
- Deeper Spark IP
- Cross-cloud support
- Single node support
- Maturer autoscale
- Tight MLFlow integration
- Tighter Delta Lake in Spark SQL
- Spark SQL Live Tables
- Component level Azure Data Factory integration
- Notebook as Dashboard

Hive Support
Java/Scala/Python
Spark Scaleout
GPU Support
SparkML Support

- SynapseML support
- Serverless or Dedicated SQL Pool
- Data warehousing
- Data Lake Tables
- Tighter Power BI integration
- .NET language support
- Big Data Scale Model Scoring with Zero Data Movement
- Transact-SQL compatibility
- Azure Purview Integration
- Azure Data Factory workspace integration




AZURE
SYNAPSE

Azure Machine Learning



- Cloud platform for creating and operating machine learning solutions



Azure Machine Learning – Deep Learning Benefits

DIVERGENCE ONE > 2022dsiwazureml > Data Labeling > Create project

- Project details
- Add workforce (optional)
- Select or create data
- Incremental refresh (optional)
- Label classes
- Labeling instructions (optional)
- ML assisted labeling (optional)

Project details

ⓘ New feature: To make labeling faster, we've added a new feature to train an ML model while you label. This feature cur

Project name *
Labeling_Sample_Project

Media type *
 Image Text

Labeling task type *

 Image Classification Multi-class	 Image Classification Multi-label	 Object Identification (Bounding Box)	 Instance Segmentation (Polygon)
---	---	---	--

Apply only a single class from a set of classes to an image

[Learn more](#)

- Labeling is a common technique to add target columns to raw data acquired prior to formal deep learning
- ML assisted = Semi-supervised learning




Azure Machine Learning – Deep Learning Benefits

The screenshot shows the Azure Machine Learning Studio interface. On the left is a vertical toolbar with various icons for data management, compute, pipelines, and monitoring. The main area has a header "DIVERGENCE ONE > 2022dsiwazureml". Below it is a "Welcome to the Azure Machine Learning Studio" message. There are four main cards: "Create new" (Notebook, Automated ML job, Pipeline), "Notebooks" (with a "Start now" button), "Automated ML" (with a "Start now" button), and "Designer" (with a "Start now" button). A "Recent" sidebar lists items like Data asset, Compute instance, Training cluster, Datastore, Data labeling project, Environment, Endpoint, and Job (preview). Below this is a "Data" section and a table of recent experiments:

	Experiment	Status	Submitted time	Submitted by	Job type
...	AdultCensusRTP	Completed	Jun 14, 2022 10:...	Vijay Koju	Pipeline
...	AdultCensusRTP	Completed	Jun 14, 2022 10:...	Vijay Koju	Pipeline
...	AdultCensusRTP	Completed	Jun 14, 2022 9:4...	Vijay Koju	Pipeline
...	AdultCensusRTP	Failed ⓘ	Jun 14, 2022 9:3...	Vijay Koju	Pipeline
Adult Census Classification	Adult-Census-Vijay	Completed	Jun 10, 2022 6:2...	Vijay Koju	Pipeline

- Scale-up
- Azure ML cluster of attached compute
- JupyterHub integration
- R Studio VM integration
- Easy integration with full ML lifecycle


Azure Synapse Analytics - Architecture



Azure Synapse Analytics – Deep Learning Benefits



- Scale-out
- Spark integration built-in
- Easier Azure DevOps & Azure Data Factory integration
- .NET language, Transact-SQL integration



<https://www.youtube.com/watch?v=jbhnZlpCu-U>

Azure Cognitive Services



Prepackaged AI services you can integrate into solutions


Language	Speech	Vision	Decision
<ul style="list-style-type: none">• Text analysis• Question answering• Language understanding• Translation	<ul style="list-style-type: none">• Speech recognition• Speech synthesis• Speech Translation• Speaker Recognition	<ul style="list-style-type: none">• Image analysis• Video analysis• Image classification• Object detection• Facial analysis• Optical character recognition	<ul style="list-style-type: none">• Anomaly detection• Content moderation• Content personalization



Azure Applied AI Services


- Form Recognizer
- Metrics Advisor
- Video Analyzer for Media
- Immersive Reader
- Bot Service
- Cognitive Search

Azure Deep Learning Platforms – Hybrid Answers



Software Development Skills

- Coding (C#, Python, Node.js, ...)
- Consuming APIs (REST or SDKs)
- DevOps (source control, CI/CD)



Conceptual AI Understanding

- Model training and inferencing
- Probability and confidence scores
- Responsible AI and ethics



Considerations for Responsible AI



Fairness



Reliability & Safety



Privacy & Security



Inclusiveness



Transparency



Accountability

Big Questions on Cognitive Services vs BYOM



DO YOU HAVE
THE RIGHT SPONSORS
AND TALENT FOR THE TASK?



WHAT LEVEL OF RISK & REWARD
ARE YOU ASSUMING BUILDING
EVERYTHING ON YOUR OWN?



DO YOU HAVE A PESSIMISTIC
ENOUGH BUDGET
EITHER WAY?



Deep Learning Deeper Dive




Deep Learning and Activation Functions



Activation function	Equation	Example	1D Graph
Unit step (Heaviside)	$\phi(z) = \begin{cases} 0, & z < 0, \\ 0.5, & z = 0, \\ 1, & z > 0, \end{cases}$	Perceptron variant	
Sign (Sigmoid)	$\phi(z) = \begin{cases} -1, & z < 0, \\ 0, & z = 0, \\ 1, & z > 0, \end{cases}$	Perceptron variant	
Linear	$\phi(z) = z$	Adaline, linear regression	
Piece-wise linear	$\phi(z) = \begin{cases} 1, & z \geq \frac{1}{2}, \\ z + \frac{1}{2}, & -\frac{1}{2} < z < \frac{1}{2}, \\ 0, & z \leq -\frac{1}{2}, \end{cases}$	Support vector machine	
Logistic (sigmoid)	$\phi(z) = \frac{1}{1 + e^{-z}}$	Logistic regression, Multi-layer NN	
Hyperbolic tangent	$\phi(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$	Multi-layer Neural Networks	
Rectifier, ReLU (Rectified Linear Unit)	$\phi(z) = \max(0, z)$	Multi-layer Neural Networks	
Rectifier, softplus	$\phi(z) = \ln(1 + e^z)$	Multi-layer Neural Networks	

Copyright © Sebastian Raschka 2016
(<http://sebastianraschka.com>)




<https://www.pinterest.com/pin/672232681856671318/>

Deep Learning General Model Structure Examples



Basic Structures

	Backfed Input Cell
	Input Cell
	Noisy Input Cell
	Hidden Cell
	Probabilistic Hidden Cell
	Spiking Hidden Cell
	Output Cell
	Match Input Output Cell
	Recurrent Cell
	Memory Cell
	Different Memory Cell
	Kernel
	Convolution or Pool



<https://www.kaggle.com/getting-started/151100>


Deep Learning General Model Structure Examples




Next Generation Complexity




Support Vector Machine (SVM)




Auto Encoder (AE)




Generative Adversarial Network (GAN)




Deep Feed Forward (DFF)



Deep Belief Network (DBN)



Markov Chain (MC)




<https://www.kaggle.com/getting-started/151100>

Deep Learning General Model Structure Examples




Sequence to Sequence (Seq2Seq)

one to one one to many many to one many to many many to many




Fixed-sized input to fixed-sized output (e.g. image classification)

one to many




Sequence output (e.g. image captioning takes an image and outputs a sentence of words).

many to one




Sequence input (e.g. sentiment analysis where a given sentence is classified as expressing positive or negative sentiment).

many to many



Sequence input and sequence output (e.g. Machine Translation: an RNN reads a sentence in English and then outputs a sentence in French)

many to many



Synced sequence input and output (e.g. video classification where we wish to label each frame of the video)






https://mett29.github.io/posts/2019/12/seq2seq_and_attention/

Deep Learning General Model Structure Examples



Next Generation Complexity




<https://www.kaggle.com/getting-started/151100>



Why Simple Neural Networks don't cut it



- We know it is good to learn a small model.
- From this fully connected model, do we really need all the edges?
- Can some of these be shared?




<https://cs.uwaterloo.ca/~mli/Deep-Learning-2017-Lecture5CNN.pdf>

Consider learning an image:




- Some patterns are much smaller than the whole image

Can represent a small region with fewer parameters



Same pattern appears in different places can be compressed!


What about training a lot of such “small” detectors and each detector must “move around”.



Convolutional Layers



A CNN is a neural network with some convolutional layers (and some other layers). A convolutional layer has a number of filters that does convolutional operation.



Convolution



These are the network parameters to be learned.

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

-1	1	-1
-1	1	-1
-1	1	-1

Filter 2

Each filter detects a : :
small pattern (3 x 3).

Convolution



1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

stride=1

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

Dot
product



3

-1

6 x 6 image

Convolution




1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

If stride=2

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0




6 x 6 image


Convolution




stride=1



6 x 6 image



Filter 1



Convolution



-1	1	-1
-1	1	-1
-1	1	-1


Filter 2

stride=1

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0


6 x 6 image

Repeat this for each filter




Two 4 x 4
images
Forming
2 x 4 x 4 matrix

Color image: RGB 3 channels





Convolution vs. Fully Connected




Fully-
connected

1	0	0	0	0	0	1
0	1	0	0	0	1	0
0	0	1	1	0	0	0
1	0	0	0	1	0	0
0	1	0	0	1	0	0
0	0	1	0	1	0	0






fewer parameters!




Only connect
to 9 inputs, not
fully
connected




Fewer parameters

Even fewer parameters



The whole CNN



Max Pooling





1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

-1	1	-1
-1	1	-1
-1	1	-1

Filter 2



Why Pooling



- Subsampling pixels will not change the object

bird



Subsampling



bird

We can subsample the pixels to make image smaller
fewer parameters to characterize the image

A CNN compresses fully connected network




- Reducing number of connections
- Shared weights on the edges
- Max pooling further reduces the complexity

<https://cs.uwaterloo.ca/~mli/Deep-Learning-2017-Lecture5CNN.ppt>

Max Pooling




1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0



6 x 6 image


New image
but smaller




Each filter
is a channel

2 x 2 image


The whole CNN




The whole CNN



Flattening





19 x 19 matrix

Black: 1

white: -1

none: 0



Neural
Network

Next move
(19 x 19
positions)

Fully-connected feedforward
network can be used

But CNN performs much better

AlphaGo's policy network




The following is quotation from their Nature article:


Note: AlphaGo does not use Max Pooling.

Neural network architecture. The input to the policy network is a $19 \times 19 \times 48$ image stack consisting of 48 feature planes. The first hidden layer zero pads the input into a 23×23 image, then convolves k filters of kernel size 5×5 with stride 1 with the input image and applies a rectifier nonlinearity. Each of the subsequent hidden layers 2 to 12 zero pads the respective previous hidden layer into a 21×21 image, then convolves k filters of kernel size 3×3 with stride 1, again followed by a rectifier nonlinearity. The final layer convolves 1 filter of kernel size 1×1 with stride 1, with a different bias for each position, and applies a softmax function. The match version of AlphaGo used $k = 192$ filters; Fig. 2b and Extended Data Table 3 additionally show the results of training with $k = 128, 256$ and 384 filters.

CNN in speech recognition



CNN in text classification



For Future Reference



- <http://www.kdnuggets.com/>
- <http://www.datasciencecentral.com>





kaggle



Starting with
Deep Learning:


DEMO with Keras





Digging deeper into
Deep Learning:

DEMO with TensorFlow






**Starting with
Deep Learning:**

DEMO with PyTorch



Starting with Deep Learning:

DEMO with ONNX





Questions?





Thank you!!!

Deck & Repo at

<https://github.com/hekaplex/DFWPythoneersIntroDeepLearning>

DIVERGENCE ACADEMY

(214) 766-3897

drew.minkin@divergence.one

